

Module-4.6-Stochastic Gradient Descent

Presented by Yasin Ceran

Table of Contents


1 Gradient Descent

2 Boosting




Introduction to Gradient Descent

- Tweak parameters iterative to minimize a cost function, $L(\theta)$
- Initialize θ_0 with random values (random initialization)
- Then update θ in steps until it converges to a local minimum

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \eta_i \frac{d}{d\theta} L(\theta^{(i)})$$




Gradient Descent Convergence



$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \eta_i \frac{d}{d\theta} L(\theta^{(i)})$$

Pick a Learning Rate



$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \eta_i \frac{d}{d\theta} L(\theta^{(i)})$$

Batch Gradient Descent

- Compute the gradient of the cost function wrt each model parameter $\theta_j \rightarrow$ Partial derivative
- $\frac{\partial MSE(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)}$
- Instead of computing partial derivatives individually, we can compute the gradient vector:

$$\Delta_{\theta} MSE(\theta) = \begin{pmatrix} \frac{\partial MSE(\theta)}{\partial \theta_0} \\ \frac{\partial MSE(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial MSE(\theta)}{\partial \theta_n} \end{pmatrix} = \frac{2}{m} \mathbf{X}^T (\mathbf{X}\theta - y)$$


- This formula involves calculations over the full training set X , at each Gradient Descent step, hence Batch Gradient Descent.
- $\theta^{(next\ step)} = \theta - \eta \Delta_{\theta} MSE(\theta)$

Batch Gradient Descent Example

```

1 eta = 0.1 # learning rate
2 n_iterations = 1000
3 n = 100
4 theta = np.random.randn(2,1) # random initialization
5 for iteration in range(n_iterations):
6     gradients = 2/n * X_b.T.dot(X_b.dot(theta) - y)
7     theta = theta - eta * gradients

```




Stochastic Gradient Descent

- picks a random instance at every step
- much faster but less regular

```

1 n_epochs = 50
2 t0, t1 = 5, 50 # learning schedule hyperparameters
3 def learning_schedule(t):
4     return t0 / (t + t1)
5 theta = np.random.randn(2,1) # random initialization
6 for epoch in range(n_epochs):
7     for i in range(m):
8         random_index = np.random.randint(m)
9         xi = X_b[random_index:random_index+1]
10        yi = y[random_index:random_index+1]
11        gradients = 2 * xi.T.dot(xi.dot(theta) - yi)
12        eta = learning_schedule(epoch * m + i)
13        theta = theta - eta * gradients

```



SGDClassifier and SGDRegressor

- you can use the SGDRegressor class, which defaults to optimizing the squared error cost function

```
1 from sklearn.linear_model import SGDRegressor
2 sgd_reg = SGDRegressor(max_iter=1000, tol=1e-3, penalty=None, eta0=0.1)
3 sgd_reg.fit(X, y.ravel())
4
5 sgd_reg.intercept_, sgd_reg.coef_
6 (array([4.24365286]), array([2.8250878]))
```




Table of Contents

1 Gradient Descent

2 Boosting


Boosting

- Boosting refers to any Ensemble method that can combine several weak learners into a strong learner.
- Train predictors sequentially, each trying to correct its predecessor.
- AdaBoost (Adaptive Boosting), Gradient Boosting, XGBoosting, ...

AdaBoost

- One way for a new predictor to correct its predecessor is to pay a bit more attention to the training instances that the predecessor underfitted.
- This results in new predictors focusing more and more on the hard cases.
- It cannot be parallelized

```
1 from sklearn.ensemble import AdaBoostClassifier
2 ada_clf = AdaBoostClassifier(
3     DecisionTreeClassifier(max_depth=1, n_estimators=200,
4         algorithm="SAMME.R", learning_rate=0.5)
5 ada_clf.fit(X_train, y_train)
```



Gradient Boosting

- Gradient Boosting works by sequentially adding predictors to an ensemble, each one correcting its predecessor
- Instead of tweaking the instance weights at every iteration like AdaBoost does, this method tries to fit the new predictor to the residual errors made by the previous predictor

$$f_1(x) \approx y$$

$$f_2(x) \approx y - f_1(x)$$

$$f_3(x) \approx y - f_1(x) - f_2(x)$$

```

1 from sklearn.tree import DecisionTreeRegressor
2 tree_reg1 = DecisionTreeRegressor(max_depth=2)
3 tree_reg1.fit(X, y)
4
5 y2 = y - tree_reg1.predict(X)
6 tree_reg2 = DecisionTreeRegressor(max_depth=2)
7 tree_reg2.fit(X, y2)
8
9 y3 = y2 - tree_reg2.predict(X)
10 tree_reg3 = DecisionTreeRegressor(max_depth=2)
11 tree_reg3.fit(X, y3)
12
13 y_pred = sum(tree.predict(X_new) for tree in (tree_reg1, tree_reg2, tree_reg3))


```

Gradient Boosting, Cont.

$$f_1(x) \approx y$$

$$f_2(x) \approx y - \gamma f_1(x)$$

$$f_3(x) \approx y - \gamma f_1(x) - \gamma f_2(x)$$



Learning rate γ , *i. e.* 0.1

Gradient Boosting is Gradient Descent

Linear regression

$$L(\mathbf{x}_i, y_i, \mathbf{w}, b) = \sum_i (y_i - \hat{y}_i)^2$$

$$= \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i - b)^2$$

optimize:

$$\min_{\mathbf{w} \in \mathbb{R}^p, b \in \mathbb{R}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i - b)^2$$

gradient descent:

$$\mathbf{w}_{j+1} = \mathbf{w}_j - \gamma \frac{\partial L(\mathbf{x}_i, y_i, \mathbf{w}, b)}{\partial \mathbf{w}}$$

Gradient Boosting

$$L(y_i, \hat{y}_i) = \sum_i (y_i - \hat{y}_i)^2$$

optimize:


$$\min_{\hat{y} \in \mathbb{R}^n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

gradient descent:


$$\hat{y}_{j+1} = \hat{y}_j - \gamma \frac{\partial L(y_i, \hat{y}_i)}{\partial \hat{y}}$$

Gradient Boosting Regressor


Residual prediction step 1




Total prediction step 1




Residual prediction step 2




Total prediction step 2




Residual prediction step 5




Total prediction step 5



Residual prediction step 9



Total prediction step 9



Summary

When to use tree-based models:

- Model non-linear relationships
- Doesn't care about scaling, no need for feature engineering
- Single tree: very interpretable (if small)
- Random forests very robust, good benchmark